



Linux for Beginners

A Practical Guide

Chapter 8 of 12

Networking Basics: Connections, IPs & SSH

Contents

8.1	How Linux Networking Works
8.2	Network Interfaces and IP Addresses
8.3	Testing Connectivity with ping
8.4	Downloading Files: curl and wget
8.5	Inspecting Open Connections: ss and netstat
8.6	DNS: How Names Become IP Addresses
8.7	SSH: Connecting to Remote Servers
8.8	SSH Keys: Passwordless and Secure Authentication
8.9	The SSH Config File
8.10	Transferring Files: scp and rsync
8.11	SSH Tunnels and Port Forwarding
8.12	Firewalls with ufw
8.13	Practical Networking Scenarios
8.14	Essential Commands: Quick Reference
8.15	Chapter Summary

Chapter 8

Networking Basics: Connections, IPs & SSH

8.1 How Linux Networking Works

Linux is built for networking. The majority of the world's servers, cloud infrastructure, and internet backbone runs on Linux — and understanding how Linux handles network connections gives you insight into how the internet itself works.

At its core, Linux networking follows the same layered model as all modern networking: physical hardware (a network card or Wi-Fi adapter) connects to a network, gets an IP address, and communicates using protocols like TCP and UDP. Linux exposes all of this through the filesystem — network interfaces appear as named devices, and everything is configurable from the command line.

Key networking concepts

Concept	What It Means
IP address	A unique numerical label identifying a device on a network (e.g. 192.168.1.5)
Subnet mask	Defines which part of an IP is the network and which is the host
Gateway	The router address — where traffic goes when leaving the local network
DNS	Domain Name System — translates names like google.com to IP addresses
Port	A numbered endpoint on a host — e.g. port 22 (SSH), 80 (HTTP), 443 (HTTPS)
Interface	A network device: eth0 (wired), wlan0 (wireless), lo (loopback)
TCP	Connection-oriented protocol — reliable, ordered delivery (web, SSH, email)
UDP	Connectionless protocol — fast but no guarantee of delivery (DNS, video)

8.2 Network Interfaces and IP Addresses

A **network interface** is how Linux represents a network connection. Each physical or virtual network adapter has an interface name. The primary tool for inspecting and configuring interfaces is the `ip` command.

```
# Show all network interfaces and their IP addresses:
$ ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500
    link/ether 52:54:00:ab:cd:ef brd ff:ff:ff:ff:ff:ff
    inet 192.168.1.100/24 brd 192.168.1.255 scope global eth0
3: wlan0: <BROADCAST,MULTICAST> mtu 1500
    link/ether aa:bb:cc:dd:ee:ff brd ff:ff:ff:ff:ff:ff

# Shorter alias:
$ ip a

# Show routing table:
$ ip route show
default via 192.168.1.1 dev eth0
192.168.1.0/24 dev eth0 proto kernel scope link src 192.168.1.100
```

Common interface names

Interface	What It Represents
lo	Loopback — always 127.0.0.1 — for local communication within the machine
eth0 / ens3	Wired Ethernet — naming varies by system (older: eth0, newer: ens3, enp2s0)
wlan0 / wlp	Wireless (Wi-Fi) interface
docker0	Virtual bridge created by Docker for container networking
tun0 / tap0	Virtual tunnel interfaces — used by VPNs

8.3 Testing Connectivity with ping

ping sends ICMP echo requests to a host and measures the round-trip time. It's the first tool to reach for when diagnosing network problems.

```
bash
# Ping a host (runs forever – Ctrl+C to stop):
$ ping google.com
PING google.com (142.250.80.46) 56(84) bytes of data.
64 bytes from 142.250.80.46: icmp_seq=1 ttl=118 time=12.4 ms
64 bytes from 142.250.80.46: icmp_seq=2 ttl=118 time=11.9 ms
^C
--- google.com ping statistics ---
2 packets transmitted, 2 received, 0% packet loss
rtt min/avg/max/mdev = 11.9/12.1/12.4/0.2 ms

# Ping only 4 times:
$ ping -c 4 google.com

# Ping an IP directly (bypasses DNS):
$ ping 8.8.8.8

# Diagnose step by step:
$ ping 127.0.0.1      # test loopback (is networking up at all?)
$ ping 192.168.1.1    # test gateway (can you reach the router?)
$ ping 8.8.8.8        # test internet (can you reach the internet?)
$ ping google.com     # test DNS (can you resolve names?)
```

8.4 Downloading Files: curl and wget

Two essential tools for fetching data from the internet directly in the terminal. Both are widely used in scripts, automation, and everyday Linux work.

curl — transfer data to/from URLs

```
# Download a file:
$ curl -O https://example.com/file.tar.gz

# Download and save with a custom name:
$ curl -o myfile.tar.gz https://example.com/file.tar.gz

# Follow redirects (important for many URLs):
$ curl -L https://example.com/redirect

# Show HTTP headers only:
$ curl -I https://example.com

# Send a POST request with JSON:
$ curl -X POST -H 'Content-Type: application/json' \
  -d '{"name": "ashik"}' https://api.example.com/users

# Test an API endpoint:
$ curl https://api.github.com/users/ashihh07
```

bash

wget — download files and mirror websites

```
# Download a file:
$ wget https://example.com/file.tar.gz

# Download quietly (no progress output):
$ wget -q https://example.com/file.tar.gz

# Continue an interrupted download:
$ wget -c https://example.com/largefile.iso

# Download in background:
$ wget -b https://example.com/largefile.iso
```

bash

Feature	curl	wget
Primary use	Transferring data (upload + download)	Downloading files
HTTP methods	Full support (GET, POST, PUT, DELETE)	GET only by default
API interaction	Excellent	Limited
Resume download	Partial support	Yes (-c flag)
Recursive download	No	Yes (--recursive)
Scripting	Very common	Common for file downloads

8.5 Inspecting Open Connections: ss and netstat

These tools show what network connections are open on your system — which ports are listening, which connections are established, and which processes own them.

```
# ss is the modern replacement for netstat:

# Show all listening ports:
$ ss -tlnp
State  Recv-Q  Send-Q  Local Address:Port  Peer Address:Port  Process
LISTEN 0        128      0.0.0.0:22          0.0.0.0:*          users:(( 'sshd',pid=423))
LISTEN 0        511      0.0.0.0:80          0.0.0.0:*          users:(( 'nginx',pid=1423))

# Show all established connections:
$ ss -tnp

# Show UDP connections:
$ ss -unp

# Show all sockets (listening + connected):
$ ss -anp

# Legacy netstat (still available on many systems):
$ netstat -tlnp
```

Flag	Meaning
-t	TCP sockets
-u	UDP sockets
-l	Listening sockets only
-n	Show numbers not names (faster)
-p	Show process that owns the socket
-a	All sockets (listening and connected)

8.6 DNS: How Names Become IP Addresses

When you type `google.com`, Linux needs to find its IP address before it can connect. This lookup is handled by **DNS** (Domain Name System). Three tools let you query and debug DNS from the terminal.

```
bash
# dig – the most detailed DNS lookup tool:
$ dig google.com
;; ANSWER SECTION:
google.com. 300 IN A 142.250.80.46

# Look up a specific record type:
$ dig google.com MX      # mail servers
$ dig google.com NS      # nameservers
$ dig google.com AAAA    # IPv6 address

# Reverse lookup – IP to hostname:
$ dig -x 8.8.8.8

# nslookup – simpler DNS query:
$ nslookup google.com
Server: 127.0.0.53
Address: 127.0.0.53#53
Name: google.com
Address: 142.250.80.46

# host – the quickest DNS check:
$ host google.com
google.com has address 142.250.80.46
google.com mail is handled by 10 smtp.google.com.

# View your DNS resolver config:
$ cat /etc/resolv.conf
```

8.7 SSH: Connecting to Remote Servers

SSH (Secure Shell) is how you connect to and control remote Linux servers. Every command you've learned in this guide can be used over SSH on a machine across the room or across the world. It's the most important networking tool in Linux.

SSH encrypts everything — your password, your commands, and any data transferred. Nothing travels in plaintext. This is what makes it safe to use over untrusted networks.

Basic SSH usage


```
bash
# Connect to a remote server:
$ ssh username@hostname
$ ssh ashik@192.168.1.50
$ ssh ashik@server.example.com

# Connect on a non-standard port (default is 22):
$ ssh -p 2222 ashik@server.example.com

# Run a single command on the remote server without opening a shell:
$ ssh ashik@server.example.com 'df -h'

# Connect with verbose output (useful for debugging):
$ ssh -v ashik@server.example.com
```

What happens when you first connect

```
bash
$ ssh ashik@192.168.1.50
The authenticity of host '192.168.1.50' can't be established.
ED25519 key fingerprint is SHA256:abc123...xyz789.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '192.168.1.50' (ED25519) to known hosts.
ashik@192.168.1.50's password:
```

The first time you connect to a server, SSH shows its fingerprint and asks you to verify it. Type yes to accept — the fingerprint is saved to `~/.ssh/known_hosts`. On future connections, SSH silently verifies the server hasn't changed. If the fingerprint changes unexpectedly, SSH will warn you — a potential sign of a man-in-the-middle attack.

8.8 SSH Keys: Passwordless and Secure Authentication

Typing a password every time you SSH is slow and less secure than using **SSH keys**. A key pair consists of a **private key** (stays on your machine, never shared) and a **public key** (placed on the server). Authentication happens cryptographically — no password is ever sent over the network.

Generating an SSH key pair

```
bash
# Generate a new ED25519 key pair (recommended – modern and fast):
$ ssh-keygen -t ed25519 -C 'ashik@mycomputer'
Generating public/private ed25519 key pair.
Enter file in which to save the key (/home/ashik/.ssh/id_ed25519): [Enter]
Enter passphrase (empty for no passphrase): [Enter or type passphrase]
Your identification has been saved in /home/ashik/.ssh/id_ed25519
Your public key has been saved in /home/ashik/.ssh/id_ed25519.pub

# View your public key (this is what you share):
$ cat ~/.ssh/id_ed25519.pub
ssh-ed25519 AAAA...longstring... ashik@mycomputer
```

Copying your public key to a server

```
bash
# The easiest way – ssh-copy-id does it automatically:
$ ssh-copy-id ashik@192.168.1.50
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed.
ashik@192.168.1.50's password:
Number of key(s) added: 1

# After this, you can SSH without a password:
$ ssh ashik@192.168.1.50
# Logs in immediately – no password prompt

# Manual method (if ssh-copy-id isn't available):
$ cat ~/.ssh/id_ed25519.pub | ssh ashik@server 'mkdir -p ~/.ssh && cat >> ~/.ssh/authorized_keys'
```

SSH key file permissions

```
bash
# SSH is strict about key file permissions:
$ chmod 700 ~/.ssh
$ chmod 600 ~/.ssh/id_ed25519      # private key – only you can read
$ chmod 644 ~/.ssh/id_ed25519.pub # public key – readable by others
$ chmod 600 ~/.ssh/authorized_keys
```

WARNING

Always set a passphrase on your private key. If someone steals your private key file without a passphrase, they have immediate access to every server it's authorised on. Use `ssh-agent` to avoid typing the passphrase repeatedly.

8.9 The SSH Config File

Typing `ssh ashik@192.168.1.50 -p 2222 -i ~/.ssh/special_key` every time is tedious. The SSH config file lets you define shortcuts for your connections.

```
$ nano ~/.ssh/config

# Add entries like this:
Host myserver
    HostName 192.168.1.50
    User ashik
    Port 2222
    IdentityFile ~/.ssh/id_ed25519

Host github
    HostName github.com
    User git
    IdentityFile ~/.ssh/github_key

Host *
    ServerAliveInterval 60
    ServerAliveCountMax 3
```

```
# Now connect with just:
$ ssh myserver
# Instead of: ssh ashik@192.168.1.50 -p 2222 -i ~/.ssh/id_ed25519

# Set correct permissions on the config file:
$ chmod 600 ~/.ssh/config
```

TIP

The `Host *` block at the bottom applies to all connections. `ServerAliveInterval 60` sends a keepalive every 60 seconds so your SSH sessions don't drop due to inactivity.

8.10 Transferring Files: `scp` and `rsync`

SSH isn't just for interactive shells — it's also the transport layer for secure file transfers.

`scp` — secure copy

```
bash
# Copy a local file TO a remote server:
$ scp file.txt ashik@server:/home/ashik/

# Copy a file FROM a remote server:
$ scp ashik@server:/home/ashik/file.txt .

# Copy a directory recursively:
$ scp -r my_project/ ashik@server:/home/ashik/

# Use a non-standard port:
$ scp -P 2222 file.txt ashik@server:/tmp/
```

rsync — efficient sync with delta transfers

rsync only transfers the parts of files that have changed. For large files or frequent syncs, it's dramatically faster than scp.

```
bash
# Sync a directory to a remote server:
$ rsync -avz my_project/ ashik@server:/home/ashik/my_project/

# Sync from remote to local:
$ rsync -avz ashik@server:/var/log/ ./server_logs/

# Dry run — show what would be transferred without doing it:
$ rsync -avzn my_project/ ashik@server:/home/ashik/my_project/

# Delete files on destination that no longer exist on source:
$ rsync -avz --delete my_project/ ashik@server:/home/ashik/my_project/
```

Flag	Meaning
-a	Archive mode — preserves permissions, timestamps, symlinks, recursion
-v	Verbose — show files being transferred
-z	Compress during transfer — faster over slow connections
-n	Dry run — simulate without making changes
--delete	Remove destination files that no longer exist in source
--progress	Show transfer progress per file

8.11 SSH Tunnels and Port Forwarding

SSH can forward network ports through an encrypted tunnel — a powerful feature for accessing services on remote machines as if they were local, or for securely exposing local services.

Local port forwarding

Makes a port on the remote server accessible as a local port on your machine:

```
# Access a remote database (port 5432) as localhost:5432:
$ ssh -L 5432:localhost:5432 ashik@server

# Syntax: -L local_port:remote_host:remote_port
# Now connect locally:
$ psql -h localhost -p 5432 -U dbuser mydb
```

bash

Remote port forwarding

Exposes a local port on the remote server — useful for sharing a local dev server:

```
# Expose your local web server (port 3000) on the remote server's port 8080:
$ ssh -R 8080:localhost:3000 ashik@server
# Syntax: -R remote_port:local_host:local_port
```

bash

8.12 Firewalls with ufw

Linux includes a powerful firewall called **iptables** (and the newer **nftables**), but configuring it directly is complex. **ufw** (Uncomplicated Firewall) is a beginner-friendly frontend that Ubuntu installs by default.

```
bash
# Check firewall status:
$ sudo ufw status
Status: inactive

# Enable the firewall:
$ sudo ufw enable

# Allow SSH (do this BEFORE enabling if connecting remotely!):
$ sudo ufw allow ssh
$ sudo ufw allow 22

# Allow common services:
$ sudo ufw allow http      # port 80
$ sudo ufw allow https    # port 443
$ sudo ufw allow 3000     # custom port

# Deny a port:
$ sudo ufw deny 8080

# Delete a rule:
$ sudo ufw delete allow 3000

# Check status with rule numbers:
$ sudo ufw status numbered

# Allow from a specific IP only:
$ sudo ufw allow from 192.168.1.50 to any port 22
```

WARNING

Always allow SSH (`sudo ufw allow ssh`) BEFORE enabling the firewall on a remote server. Enabling `ufw` without an SSH rule will lock you out immediately.

8.13 Practical Networking Scenarios

Scenario 1: Diagnose no internet connection

```
bash
# Step 1: is the interface up?
$ ip addr show eth0

# Step 2: can you reach the router?
$ ping -c 3 192.168.1.1

# Step 3: can you reach the internet by IP?
$ ping -c 3 8.8.8.8

# Step 4: is DNS working?
$ ping -c 3 google.com
$ dig google.com
```

Scenario 2: Set up SSH key-based login on a new server

```
bash
# On your local machine:
$ ssh-keygen -t ed25519 -C 'mykey'
$ ssh-copy-id user@newserver

# Verify it works:
$ ssh user@newserver

# On the server – optionally disable password auth:
$ sudo nano /etc/ssh/sshd_config
# Set: PasswordAuthentication no
$ sudo systemctl restart sshd
```

Scenario 3: Check what's listening on a port

```
bash
# Is anything listening on port 80?
$ ss -tlnp | grep :80
LISTEN 0 511 0.0.0.0:80 0.0.0.0:* users:(("nginx",pid=1423,fd=6))

# What process is using port 3000?
$ ss -tlnp | grep :3000
```

Scenario 4: Sync a local project to a server

```
bash
# Dry run first to check what will be sent:
$ rsync -avzn ~/my_project/ ashik@server:/var/www/my_project/

# If the dry run looks correct, run for real:
$ rsync -avz ~/my_project/ ashik@server:/var/www/my_project/
```

8.14 Essential Commands: Quick Reference

Command	What It Does	Common Usage
<code>ip addr show</code>	Show network interfaces and IPs	<code>ip a</code>
<code>ip route show</code>	Show routing table and gateway	<code>ip route show</code>
<code>ping host</code>	Test connectivity to a host	<code>ping -c 4 google.com</code>
<code>curl -O URL</code>	Download a file	<code>curl -O https://...</code>
<code>curl -I URL</code>	Show HTTP response headers only	<code>curl -I https://example.com</code>
<code>wget URL</code>	Download a file	<code>wget https://...</code>
<code>ss -tlnp</code>	Show listening TCP ports and processes	<code>ss -tlnp</code>
<code>ss -anp</code>	Show all sockets	<code>ss -anp</code>
<code>dig domain</code>	DNS lookup with full details	<code>dig google.com MX</code>
<code>nslookup domain</code>	Simple DNS lookup	<code>nslookup google.com</code>
<code>host domain</code>	Quick DNS lookup	<code>host google.com</code>
<code>ssh user@host</code>	Connect to a remote server	<code>ssh ashik@192.168.1.50</code>
<code>ssh -p PORT user@host</code>	SSH on a non-standard port	<code>ssh -p 2222 ashik@server</code>
<code>ssh-keygen -t ed25519</code>	Generate an SSH key pair	<code>ssh-keygen -t ed25519</code>
<code>ssh-copy-id user@host</code>	Copy public key to server	<code>ssh-copy-id ashik@server</code>
<code>scp file user@host:path</code>	Copy file to remote server	<code>scp app.py ashik@server:~/</code>
<code>scp user@host:path .</code>	Copy file from remote server	<code>scp ashik@server:~/log.txt .</code>
<code>rsync -avz src dst</code>	Sync files efficiently over SSH	<code>rsync -avz project/ server:~/</code>
<code>ssh -L lport:host:rport</code>	Local port forwarding tunnel	<code>ssh -L 5432:localhost:5432 server</code>
<code>sudo ufw status</code>	Check firewall status	<code>sudo ufw status</code>
<code>sudo ufw allow ssh</code>	Allow SSH through firewall	<code>sudo ufw allow ssh</code>
<code>sudo ufw enable</code>	Enable the firewall	<code>sudo ufw enable</code>

8.15 Chapter Summary

- ✓ Linux networking is built around **interfaces** (eth0, wlan0, lo), **IP addresses**, **routes**, and **DNS**. The `ip` command is the primary tool for inspecting and configuring all of these.
- ✓ `ping` is the first diagnostic tool — test loopback, then gateway, then internet IP, then DNS to isolate any connectivity problem.
- ✓ `curl` transfers data and interacts with APIs. `wget` downloads files and supports resuming interrupted downloads.
- ✓ `ss -tlnp` shows what's listening on which ports and which process owns each socket. Use it to confirm services are running and to find port conflicts.
- ✓ SSH is the foundation of remote Linux administration — encrypted, reliable, and available everywhere. Learn it thoroughly.
- ✓ **SSH keys** are more secure than passwords and faster to use. Always use ED25519, always set a passphrase, always set correct permissions (`chmod 600` on private keys).
- ✓ The `~/.ssh/config` file lets you define shortcuts for all your SSH connections — host aliases, ports, users, and key files in one place.
- ✓ `scp` copies files over SSH simply. `rsync` syncs directories efficiently — only transferring what changed. Always do a dry run (`-n`) before a real `rsync`.
- ✓ SSH tunnels forward ports through encrypted connections — useful for accessing remote databases, services, and APIs as if they were local.
- ✓ `ufw` makes firewall management approachable. Always allow SSH before enabling the firewall on a remote server.

What's Coming in Chapter 9

With networking mastered, Chapter 9 brings everything together: **shell scripting**. You'll learn to automate tasks by combining everything you've learned into reusable scripts:

- What a shell script is and how the shebang line (`#!/bin/bash`) works
- Variables, user input, and command substitution
- Conditionals: `if`, `elif`, `else`, and test expressions
- Loops: `for`, `while`, and `until`
- Functions — reusable blocks of logic
- Reading arguments, handling errors, and exit codes
- Practical scripts: backups, log rotation, system health checks

Chapter 9 is where Linux stops being a tool you use and starts being a tool you program.

Linux for Beginners — Chapter 8 of 12

Author

Ashik A

github.com/ashihh07

This guide was created, organized, refined, and designed by Ashik A using GitHub documentation, industry best practices, publicly available learning resources, and AI-assisted research.
